

Frequently Asked Questions (FAQ) - AI Features

Applicable Product: All AI-Enhanced NexaCore Modules (PAWT, Summarization, Classification)

Audience: Enterprise Users

Date: 2025-11-01

1. Data Security and Privacy

Q: Is my data used to train the global NexaCore AI models?

A: No. Your client data is only used to *fine-tune* your specific private instance of our model (tenant-level isolation). Data is securely encrypted and stored according to our strict regional data residency policies. We do not use client data to train general, public models.

Q: Where is my data physically stored?

A: If you are an Enterprise client in the MENA region, your data is stored on regional servers located in the UAE/KSA, as required by your service agreement. This is guaranteed by our data residency compliance policy.

2. AI Model Performance

Q: My AI model is giving inaccurate predictions. What should I do?

A: This is usually due to **Data Drift**. The model was trained on historical data, but your operational patterns have changed (e.g., a major seasonal shift in volume).

1. Check the **Model Health Dashboard** for a "Low Confidence" warning.
2. If confirmed, submit a P2 ticket with the model ID and request a **Re-Training** with the latest 6 months of your data.

Q: How often is the model retrained?

A: By default, all enterprise models are set for an automated quarterly re-training cycle. However, you can request an ad-hoc re-train if you notice sustained performance degradation.

3. Billing and Usage

Q: Why did my bill increase after activating the Predictive AI Workflow Trigger (PAWT)?

A: The PAWT feature is part of our Premium Subscription Tier. Billing is based on the **number of inference calls** the model makes per month, not the number of successful workflow

triggers. You can monitor your current call volume in the **Billing Portal** dashboard.